

COMP462

Project 2: Robotic Grasping

In this project, we will evaluate the quality of contact-based grasps, and then optimize the contacts to improve the grasp quality. A code framework, hosted on github <https://github.com/robotpilab/COMP462/tree/main/project2>, is provided to facilitate the implementation.

Report: A report is required for this project. The requirements of what to be included in the report are highlighted in the description.

The object mesh models, which are the grasping targets, are provided and managed by the `trimesh` library in the provided code framework. Essential properties of the mesh models can be easily accessed, including the number of faces/vertices, the coordinates of vertices, the normals of faces, face adjacencies, etc. Please refer to the library documentation (<https://trimesh.org/trimesh.html>) and examples (https://trimesh.org/examples.quick_start.html) for a quick start.

A grasp on an object is represented by the indices of the faces (i.e., triangles) where the contacts are. For example, if a grasp is defined by `grasp=[2, 7, 36, 94]`, this means the object is grasped with contacts at its 2-nd, 7-th, 36-th, and 94-th faces. The faces of an object provided in this project are all triangles. By convention, we use the centroids of the triangular faces as the coordinates of the contact points.

To calculate the centroids given the indices of triangles, the following function is provided in `utils.py`:

```
def get_centroid_of_triangles(mesh, tr_ids):
    """
    Calculate the centroids of the triangles.
    args:  mesh: The object mesh model.
            Type: trimesh.base.Trimesh
          tr_ids: The indices of the triangles on the mesh model.
            Type: list of int
    returns: cen: The centroids of the triangles.
              Type: numpy.ndarray of shape (len(tr_ids), 3)
    """
```

Task 1: Grasp Quality Evaluation

Given an object mesh model and a grasp, calculate the L_1 quality score of this grasp. Major steps for this task are detailed below:

1. Obtain the primitive wrenches for a grasp given the friction coefficient. Specifically, you need to implement the following function in `alg.py`:

```
def primitive_wrenches(mesh, grasp, mu=0.2, n_edges=8):
    """
    Find the primitive wrenches for each contact of a grasp.
    args:  mesh: The object mesh model.
            Type: trimesh.base.Trimesh
          grasp: The indices of the mesh triangles being contacted.
            Type: list of int
          mu: The friction coefficient of the mesh surface.
              (default: 0.2)
          n_edges: The number of edges of the friction polyhedral cone.
            Type: int (default: 8)
```

```

returns: W: The primitive wrenches.
          Type: numpy.ndarray of shape (len(grasp) * n_edges, 6)
"""

```

Given the friction coefficient μ of the object surface, the first step is to compute the polyhedral friction cone for each contact, which is represented by `n_edges` number of primitive forces. Then for each primitive force \mathbf{f} , compute the corresponding torque $\boldsymbol{\tau}$ generated by this force to obtain the primitive wrench. In the end, the returned value `W` of the function should be a `numpy.ndarray`, each row of which is one primitive wrench $[\mathbf{f}^\top, \boldsymbol{\tau}^\top]$.

For example, given a grasp `grasp = [2, 7, 36, 94]` consisting of four contacts, the returned `W` should be of shape `(4 * n_edges, 6)`. Every `n_edges` rows of `W` are the primitive wrenches for one contact of the grasp.

2. Next, you need to implement the following function in `alg.py` to evaluate the L_1 quality of a given grasp, based on the primitive wrenches you implemented in the previous step.

```

def eval_Q(mesh, grasp, mu=0.2, n_edges=8, lmbd=1.0):
    """
    Evaluate the L1 quality of a grasp.
    args:  mesh: The object mesh model.
            Type: trimesh.base.Trimesh
          grasp: The indices of the mesh triangles being contacted.
            Type: list of int
          mu: The friction coefficient of the mesh surface.
            (default: 0.2)
          n_edges: The number of edges of the friction polyhedral cone.
            Type: int (default: 8)
          lmbd: The scale of torque magnitude.
            (default: 1.0)
    returns: Q: The L1 quality score of the given grasp.
    """

```

As introduced in the lecture, the L_1 quality is defined as the distance from the origin to the nearest hyperplane of the convex hull spanned by primitive wrenches. You can import and use the `scipy.spatial` package for convex hull-related calculations.

The distance of the wrench space can be defined in different ways by scaling the torque magnitude with a parameter λ (1.0 by default in the provided code framework):

$$\|\mathbf{w}\| = \sqrt{\|\mathbf{f}\|^2 + \lambda\|\boldsymbol{\tau}\|^2}$$

After completing this part, you can run the following command in your terminal to test your implementation: `python main.py --task 1`. This will visualize the mesh model and the given grasp as exemplified in Figure 1, and print your results in the terminal.

Moreover, you can try with different mesh models `bunny`, `cow`, and `duck`. To do this, you can just add `--mesh the_model_you_choose` to the end of your command.

Report: For each one of the mesh models (`bunny`, `cow`, and `duck`), given a three-contact grasp `grasp=[0, 249, 484]`, calculate the L_1 quality when $\lambda = 0.2, 0.5$, and 1 respectively. In the report, please show a visualization of each grasp, and provide a table summarizing the grasp qualities.




Figure 1: The visualization of an example grasp on the bunny. The red dots are the contact points of this grasp and the green arrows are the normals (defined in the outward direction for better visualization) at the contacts.

Task 2: Sample a Stable Grasp

For this part, you need to randomly sample a stable grasp on a provided mesh model, that is, the quality of this grasp is above a threshold. You can call your functions implemented in Task 1. Specifically, you need to implement the following function in `grasp.py`:

```
def sample_stable_grasp(mesh, thresh=0.0):
    """
    Sample a stable grasp such that its L1 quality is larger than a threshold.
    args:     mesh: The object mesh model.
                Type: trimesh.base.Trimesh
            thresh: The threshold for stable grasp.
                (default: 0.0)
    returns: grasp: The stable grasp represented by indices of triangles.
                Type: list of int
            Q: The L1 quality score of the sampled grasp,
                expected to be larger than thresh.
    """
```

After completing this part, you can invoke the following command to test your implementation. This will visualize the mesh model and your sampled grasp and prints your results in the terminal:
`python main.py --task 2`

Report: Set $\lambda = 1.0$ and `thresh=0.01`, please provide five different grasps (each grasp has three contacts) for the bunny model. In the report, please list the grasp contacts, visualize the grasps, and provide statistical analysis of the qualities of sampled grasps, e.g., histograms.

Task 3: Grasp Optimization

Given a mesh model and an initial grasp on it, now we focus on optimizing the grasp. We will record the trajectory of the optimization steps. The trajectory should be a list of intermediate grasps during the optimization process. Below are the major steps needed for this task:

1. Implement the following function in `alg.py` to find the η -order neighbors of any face indexed by `tr_id` on the mesh model. For this, you might be interested in accessing the `face_neighborhood` property of the mesh model to retrieve the pairs of faces that share the same vertex.

```
def find_neighbors(mesh, tr_id, eta=1):
    """
    Find the eta-order neighbor faces (triangles) of tr_id on the mesh model.
    args:      mesh: The object mesh model.
                  Type: trimesh.base.Trimesh
              r_id: The index of the query face (triangle).
                  Type: int
              eta: The maximum order of the neighbor faces:
                  Type: int
    returns: nbr_ids: The list of the indices of the neighbor faces.
                  Type: list of int
    """
```

2. Implement the following function in `alg.py` to find the optimal neighbor grasp of an input grasp.

```
def local_optimal(mesh, grasp):
    """
    Find the optimal neighbor grasp of the given grasp.
    args:      mesh: The object mesh model.
                  Type: trimesh.base.Trimesh
              grasp: The indices of the mesh triangles being contacted.
                  Type: list of int
    returns: G_opt: The optimal neighbor grasp with the highest quality.
                  Type: list of int
              Q_max: The L1 quality score of G_opt.
    """
```

Two grasps are considered neighbors if their corresponding contacts are neighbored faces (triangles) on the mesh model. For this step, you might want to use your `find_neighbors()` function implemented in the previous step to find neighbor faces for each contact of the grasp. And a neighbor grasp of the input grasp is just one of the combinations of the neighbor faces,

3. Finish the following function in `alg.py` to iteratively find the optimal neighbor grasp until the grasp quality cannot be improved within its neighbors. At each iteration, remember to trace the locally optimized grasp as you need to return the whole trajectory in the end.

```
def optimize_grasp(mesh, grasp):
    """
    Optimize the given grasp and return the trajectory.
    args:      mesh: The object mesh model.
                  Type: trimesh.base.Trimesh
              grasp: The indices of the mesh triangles being contacted.
                  Type: list of int
    returns: traj: The trajectory of the grasp optimization.
                  Type: list of grasp (each grasp is a list of int)
    """
```

After completing this part, you can run the following command in the terminal to test your implementation. This will visualize the mesh model and your trajectory, as exemplified in Figure 2:

```
python main.py --task 3
```




Figure 2: An example grasp optimization on the bunny. The red dots and yellow dots are the initial and optimized grasp respectively, and the green lines are the grasp trajectory during optimization.

Report: On the bunny model, with $\lambda = 1.0$, please optimize the grasp=[80, 165, 444] and visualize the optimization trajectory in the report. Please discuss: how many grasp quality evaluations were needed in this optimization process? If you randomly sample grasps, what is the probability of obtaining a grasp of your optimized quality?

Task 4: Grasp Optimization with Reachability Constraint (Optional)

In this task, given a mesh model you need to first sample a grasp under the reachability constraint, and then optimize this grasp while complying with the reachability constraint. The reachability is used to constrain the contact points of a grasp not too far away from each other. Formally, this constraint is satisfied when the average of the distances from contact points to their centroid are less than a reachability measure r :

$$\frac{1}{m} \sum_{i=1}^m \|p_i - \psi\| < r, \quad \text{where } p_1, \dots, p_m \text{ are contact points and } \psi = \frac{1}{m} \sum_{i=1}^m p_i$$

Specifically, you need to implement the following function located in `grasp.py`. Feel free to implement more functions if needed.

```
def optimize_reachable_grasp(mesh, r=0.5):
    """
    Sample a reachable grasp and optimize it.
    args:    mesh: The object mesh model.
              Type: trimesh.base.Trimesh
              r: The reachability measure. (default: 0.5)
    returns: traj: The trajectory of the grasp optimization.
              Type: list of grasp (each grasp is a list of int)
    """
```

Hint: After you sample a reachable grasp, you can use the same algorithm you implemented in Task 3 to optimize the grasp while only accepting the the reachable neighbors at each iteration.

After completing this part, you can run the following command in your terminal to test your program, which visualizes the mesh model and your trajectory, as exemplified in Figure 3: `python main.py --task 4`

Report: On the bunny model, with $r = 0.5$ and 1.0 respectively, please first sample a reachable grasp and optimize this grasp under the reachability constraint. Feel free to test with different initially sampled grasps. For each r , select one optimized trajectory to visualize in the report. Please provide a pseudo-code of your constraint-based optimization algorithm, and discuss potential improvements you can add to your algorithm.




Figure 3: An example grasp optimization with a reachability constraint of $r = 1.0$.